

# Caching or pre-fetching? The role of hazard rates.

**Andres Ferragut**

joint work with Matias Carrasco and Fernando Paganini

Universidad ORT Uruguay

60th Allerton Conference on Computing, Control and Communications – September 2024

Local memory systems

Timer based caching


Timer based pre-fetching

Asymptotic optimality

Final remarks

# The caching problem

- Consider a **local memory system** that handles items from a catalog of  $N$  objects.
- Requests for objects arrive as a random process.
- The memory (cache) can locally store  $C < N$  of them.
- If item is in cache, we have a **hit**. Otherwise, it is a **miss**.




**Objective:** for a given arrival stream, maximize the steady-state **hit rate**.

# Point process approach


Introduced in [Fofack et al. 2014]

- Assume requests for item  $i$  come from a **point process** of intensity  $\lambda_i$  (popularities).



- At each point in time we must decide which items must be stored locally.


## Two important distributions:



- **Inter-arrival distribution:** Typical distance between points...

$$\tau_{k+1} - \tau_k \sim F(t), \quad E[\tau_{k+1} - \tau_k] = 1/\lambda$$

## Two important distributions:




- **Inter-arrival distribution:** Typical distance between points...

$$\tau_{k+1} - \tau_k \sim F(t), \quad E[\tau_{k+1} - \tau_k] = 1/\lambda$$

- **Age distribution:** Distance from the last point in the current interval (sampling bias)!

$$F(t) := \lambda \int_0^t 1 - F_0(s) ds,$$

## Two important distributions:



- **Inter-arrival distribution:** Typical distance between points...

$$\tau_{k+1} - \tau_k \sim F(t), \quad E[\tau_{k+1} - \tau_k] = 1/\lambda$$


- **Age distribution:** Distance from the last point in the current interval (sampling bias)!

$$F(t) := \lambda \int_0^t 1 - F_0(s) ds,$$

**Note:** you can formalize this under the **Palm probability** framework for stationary point processes.

## Populating a cache: timer based (TTL) policies

- Upon request arrival for item  $i$ , check for presence.






## Populating a cache: timer based (TTL) policies

- Upon request arrival for item  $i$ , check for presence.
- If new, store item and start a **timer**  $T_i$  to evict.




## Populating a cache: timer based (TTL) policies

- Upon request arrival for item  $i$ , check for presence.
- If new, store item and start a **timer**  $T_i$  to evict.
- If present, reset timer to  $T_i$ .




## Populating a cache: timer based (TTL) policies

- Upon request arrival for item  $i$ , check for presence.
- If new, store item and start a **timer**  $T_i$  to evict.
- If present, reset timer to  $T_i$ .
- Upon timer expiration, evict the content.



## Populating a cache: timer based (TTL) policies

- Upon request arrival for item  $i$ , check for presence.
- If new, store item and start a **timer**  $T_i$  to evict.
- If present, reset timer to  $T_i$ .
- Upon timer expiration, evict the content.
- Keep timers  $T_i$  such that **average** cache occupation is  $C$ .



# Structure of the optimal caching policy

- The crucial magnitude is the **hazard rate** of  $F_0$ :

$$\eta(t) := \frac{f_0(t)}{1 - F_0(t)}$$

- Likelihood of a request at time  $t$ , given the current interval has age  $t$ .

# Structure of the optimal caching policy

- The crucial magnitude is the **hazard rate** of  $F_0$ :

$$\eta(t) := \frac{f_0(t)}{1 - F_0(t)}$$

- Likelihood of a request at time  $t$ , given the current interval has age  $t$ .

## Theorem (F', Rodriguez, Paganini – 2016)

If the  $F_0^{(i)}$  have **decreasing hazard rates**, then the optimal TTL policy satisfies:


$$\eta_i(T_i^*) \geq \theta^*,$$

whenever  $T_i^* > 0$  (i.e. the item is cached).

Moreover, inequality is strict iff  $T_i^* = \infty$  (item always stored).


# Why caching helps in this case?

Decreasing hazard rates corresponds to **bursty traffic**:



# Why caching helps in this case?

Decreasing hazard rates corresponds to **bursty traffic**:




- An arrival makes a subsequent arrival **more likely**.
- Store it while its likelihood is high enough (above a threshold).




# Why caching helps in this case?

Decreasing hazard rates corresponds to **bursty traffic**:




- An arrival makes a subsequent arrival **more likely**.
- Store it while its likelihood is high enough (above a threshold).

# What about other types of traffic?




Constant hazard rate  $\rightarrow$  Poisson process.




Increasing hazard rate  $\rightarrow$  more periodic!

# What about other types of traffic?




Constant hazard rate  $\rightarrow$  Poisson process.




Increasing hazard rate  $\rightarrow$  more periodic!

**Theorem:** for these types of traffic, keep the most popular is the optimal **caching** policy.

# What about other types of traffic?



Constant hazard rate  $\rightarrow$  Poisson process.




Increasing hazard rate  $\rightarrow$  more periodic!

**Theorem:** for these types of traffic, keep the most popular is the optimal **caching** policy.

Can we improve upon this?

## Thinking about increasing hazard rates...


- Once you have seen a request, it's **less likely** to see the same item again for a while.



- What is the timer based equivalent of this case?

## Thinking about increasing hazard rates...

- Once you have seen a request, it's **less likely** to see the same item again for a while.




- What is the timer based equivalent of this case?

### Key insight

The question now is not **how long we should remember something**, but instead **how long we should forget about it!**

## Timer based pre-fetching policy

- Upon request arrival for item  $i$ , check for presence.



## Timer based pre-fetching policy


- Upon request arrival for item  $i$ , check for presence.
- If not-present: start a **timer**  $T_i$ .






# Timer based pre-fetching policy

- Upon request arrival for item  $i$ , check for presence.
- If not-present: start a **timer**  $T_i$ .
- If present: remove content and reset timer to  $T_i$ .




## Timer based pre-fetching policy

- Upon request arrival for item  $i$ , check for presence.
- If not-present: start a **timer**  $T_i$ .
- If present: remove content and reset timer to  $T_i$ .
- Upon timer expiration, **pre-fetch** the content.



## Timer based pre-fetching policy


- Upon request arrival for item  $i$ , check for presence.
- If not-present: start a **timer**  $T_i$ .
- If present: remove content and reset timer to  $T_i$ .
- Upon timer expiration, **pre-fetch** the content.
- Keep timers  $T_i$  such that **average** cache occupation is  $C$ .



# Timer based pre-fetching


Consider a single item with a timer  $T$  and its request process:

**Hit probability:** next arrival occurs **after** timer expires.



$$\text{Hit probability} = 1 - F_0(T)$$

**Occupation probability:** probability that timer **has** expired by 0 since last arrival.



$$\text{Avg. occupation} = 1 - F(T)$$

# Choosing the optimal timers

Requests come from independent sources with intensities  $\lambda_i$  and inter-arrival distribution  $F_i$ :

## Problem (Optimal pre-fetching policy)

Choose timers  $T_i \geq 0$  such that:

$$\max_{T_i \geq 0} \sum_i \lambda_i (1 - F_i(T_i))$$

subject to:

$$\sum_i (1 - \hat{F}_i(T_i)) \leq C$$

# Choosing the optimal timers

Requests come from independent sources with intensities  $\lambda_i$  and inter-arrival distribution  $F_i$ :

## Problem (Optimal pre-fetching policy)

Choose timers  $T_i \geq 0$  such that:

$$\min_{T_i \geq 0} \sum_i \lambda_i F_0^{(i)}(T_i)$$

subject to:

$$\sum_i F^{(i)}(T_i) \geq N - C$$

# Choosing the optimal timers

## Change of variables

- Apply the change of variables  $u_i = F^{(i)}(T_i)$ .
- Note that  $u_i$  is the probability of **not being stored**.
- The problem becomes:

$$\min_{u_i \in [0,1]} \sum_i \lambda_i \left[ F_0^{(i)} \circ (F^{(i)})^{-1} \right] (u_i)$$

subject to:

$$\sum_i u_i \geq N - C$$

# Choosing the optimal timers

## Lagrangian duality

### ■ Objective gradient:

$$\frac{\partial}{\partial u_i} \lambda_i F^{(i)} \circ (F^{(i)})^{-1}(u_i) = \frac{\lambda_i f_0^{(i)}((F^{(i)})^{-1}(u_i))}{\lambda_i \left(1 - F_0^{(i)}((F^{(i)})^{-1}(u_i))\right)} = \eta_i((F^{(i)})^{-1}(u_i))$$

### ■ Increasing! → Proper convex optimization problem.



# Choosing the optimal timers

## Lagrangian duality

### ■ Objective gradient:

$$\frac{\partial}{\partial u_i} \lambda_i F^{(i)} \circ (F^{(i)})^{-1}(u_i) = \frac{\lambda_i f_0^{(i)}((F^{(i)})^{-1}(u_i))}{\lambda_i \left(1 - F_0^{(i)}((F^{(i)})^{-1}(u_i))\right)} = \eta_i((F^{(i)})^{-1}(u_i))$$

### ■ Increasing! → Proper convex optimization problem.

### ■ Lagrangian duality:

$$\begin{aligned} \mathcal{L}(u, \theta) &= \sum_{i=1}^N \lambda_i F_0^{(i)} \left( (F^{(i)})^{-1}(u_i) \right) + \theta \left( N - C - \sum_{i=1}^N u_i \right) \\ &= \sum_{i=1}^N \left[ \lambda_i F_0^{(i)} \left( (F^{(i)})^{-1}(u_i) \right) - \theta u_i \right] + \theta(N - C). \end{aligned}$$

## Theorem

If the  $F_0^{(i)}$  satisfy the IHR property, there exists a unique threshold  $\theta^* \geq 0$  such that the optimal timers satisfy:

$$\eta_i(T_i^*) \geq \theta^*,$$

whenever  $T_i^* < \infty$  (pre-fetching).

The inequality is strict if and only if  $T_i^* = 0$ , i.e. the content is always stored.

## Theorem

If the  $F_0^{(i)}$  satisfy the IHR property, there exists a unique threshold  $\theta^* \geq 0$  such that the optimal timers satisfy:

$$\eta_i(T_i^*) \geq \theta^*,$$

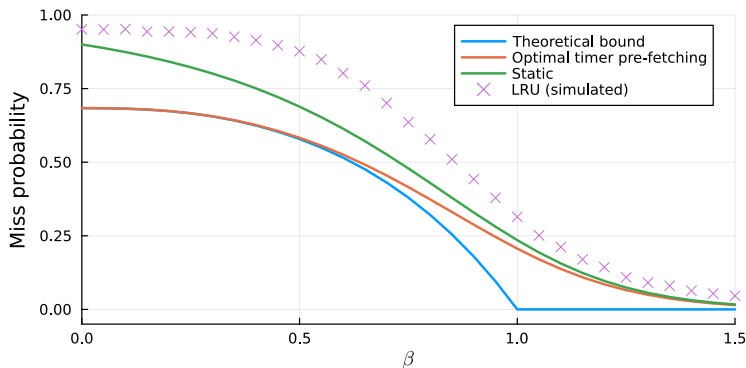
whenever  $T_i^* < \infty$  (pre-fetching).

The inequality is strict if and only if  $T_i^* = 0$ , i.e. the content is always stored.

**Remark:** The policy is also a threshold policy, like the caching case.



# Simulation example

Erlang ( $k = 5$ ) inter-arrival times, Zipf  $\propto n^{-\beta}$  popularities, varying  $\beta$ ,  $N = 10000$ ,  $C = 1000$ .





- Pre-fetching improves over the static policy.
- Classical caching (e.g. LRU) is a very bad idea for regular traffic.

# A tale of two policies...



## A tale of two policies...



Both policies are just the same policy!

- Keep a hazard rate threshold  $\theta$  for storing a content
- Compute  $\theta^*$  such that avg. memory occupation is  $C$ .

- Threshold policies in fact are related to recent results by [Panigrahy et al. 2022], about the optimal causal policy.
- They also identify a notion of hazard rate, the stochastic intensity, as defining the optimal policy.

- Threshold policies in fact are related to recent results by [Panigrahy et al. 2022], about the optimal causal policy.
- They also identify a notion of hazard rate, the stochastic intensity, as defining the optimal policy.

## Theorem (F', Carrasco, Paganini – In preparation – check ArXiv soon)

Under an appropriate scaling regime, as  $N \rightarrow \infty$  and  $C = cN$ , the optimal causal policy converges to a **fixed threshold policy**. Moreover, the limit threshold  $\theta^*$  coincides with its timer based counterpart.

Therefore, the above policies give a **universal asymptotic upper bound** on caching performance.



## Final remarks

- You have to **know your traffic** before deciding on caching or pre-fetching!

## Final remarks

- You have to **know your traffic** before deciding on caching or pre-fetching!
- Classical caching is **not well suited** to regular traffic.

- You have to **know your traffic** before deciding on caching or pre-fetching!
- Classical caching is **not well suited** to regular traffic.
- We identified the **hazard rate** as the crucial indicator of regularity, and devised a new policy for IHR, that is also **asymptotically optimal** among all causal policies.

- You have to **know your traffic** before deciding on caching or pre-fetching!
- Classical caching is **not well suited** to regular traffic.
- We identified the **hazard rate** as the crucial indicator of regularity, and devised a new policy for IHR, that is also **asymptotically optimal** among all causal policies.
- A lot of open questions, in particular:
  - How we can learn the hazard rates online?
  - How we can estimate the appropriate threshold?
  - What about mixtures of IHR and DHR traffic?

# Thank you!

Andres Ferragut

[ferragut@ort.edu.uy](mailto:ferragut@ort.edu.uy)

[aferragu.github.io](https://aferragu.github.io)

# References I



**P. Brémaud.**

Point process calculus in time and space.

Springer, NY, 2020.



**H. Che, Y. Tung, and Z. Wang.**

Hierarchical web caching systems: Modeling, design and experimental results.

IEEE Journal on Selected Areas in Communications, 20(7):1305–1314, 2002.



**A. Ferragut, I. Rodriguez, and F. Paganini.**

Optimizing TTL caches under heavy tailed demands.

In Proc. of ACM/SIGMETRICS 2016, pages 101–112, June 2016.





**A. Ferragut, I. Rodríguez, and F. Paganini.**

Optimal timer-based caching policies for general arrival processes.

Queueing Systems, 88(3–4):207–241, 2018.

## References II

-  **N. C. Fofack, P. Nain, G. Neglia, and D. Towsley.**  
Performance evaluation of hierarchical TTL-based cache networks.  
*Computer Networks*, 65:212–231, 2014.
-  **N. K. Panigrahy, P. Nain, G. Neglia, and D. Towsley.**  
A new upper bound on cache hit probability for non-anticipative caching policies.  
*ACM Trans. Model. Perform. Eval. Comput. Syst.*, 7(2–4), November 2022.